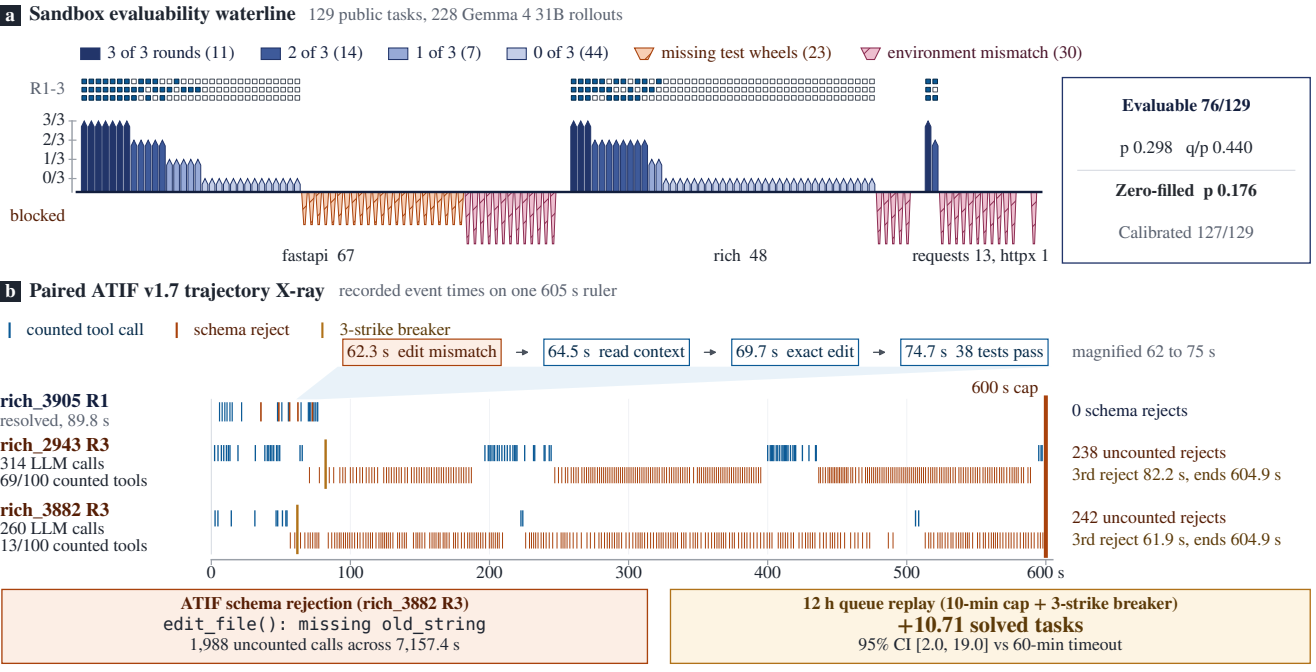
a Sandbox evaluability waterline 129 public tasks, 228 Gemma 4 31B rollouts
3 of 3 rounds (11)
2 of 3 (14)
1 of 3 (7)
0 of 3 (44)
missing test wheels (23)
environment mismatch (30)
R1-3
3/3
2/3
1/3
0/3
blocked
fastapi 67
rich 48
requests 13, httpx 1
Evaluable 76/129
p 0.298 q/p 0.440
Zero-filled p 0.176
Calibrated 127/129
b Paired ATIF v1.7 trajectory X-ray recorded event times on one 605 s ruler
counted tool call
schema reject
3-strike breaker
62.3 s edit mismatch
64.5 s read context
69.7 s exact edit
74.7 s 38 tests pass
magnified 62 to 75 s
600 s cap
rich_3905 R1
resolved, 89.8 s
0 schema rejects
rich_2943 R3
314 LLM calls
69/100 counted tools
238 uncounted rejects
3rd reject 82.2 s, ends 604.9 s
rich_3882 R3
260 LLM calls
13/100 counted tools
242 uncounted rejects
3rd reject 61.9 s, ends 604.9 s
0
100
200
300
400
500
600 s
ATIF schema rejection (rich_3882 R3)
edit_file(): missing old_string
1,988 uncounted calls across 7,157.4 s
12 h queue replay (10-min cap + 3-strike breaker)
+10.71 solved tasks
95% CI [2.0, 19.0] vs 60-min timeout